


CSCI 544: Applied Natural Language Processing

Transformer-II


Xuezhe Ma (Max)



USC University of
Southern California


Course Organization

NLP Tasks



Transformers



Transformer Encoder Block



- **Three Key Components**

- (Masked) Multi-head Self-Attention
- Layer Normalization
- Position-wise Feed-Forward Network

Transformers



Outline

- **Residual Connections**
- **Normalization Layers**
 - Variants of Normalization Layers
 - Pre-Norm vs. Post-Norm
- **Positional Embeddings**
 - Absolute Positional Embeddings
 - Rotary Positional Embeddings
- **Other Variant Components**
 - FFN -> SwiGLU
 - MQA & GQA

Residual Connections

Residual Connections

Deep Residual Learning for Image Recognition

Kaiming He Xiangyu Zhang Shaoqing Ren Jian Sun

Microsoft Research

{kahe, v-xiangz, v-shren, jiansun}@microsoft.com

Deep residual learning for image recognition

[K He](#), [X Zhang](#), [S Ren](#), [J Sun](#) - ... and [pattern recognition](#), 2016 - [openaccess.thecvf.com](#)

... **Deeper** neural **networks** are more difficult to train. We present a **residual learning** framework to ease the training of **networks** that are substantially **deeper** than those used previously. ...

☆ Save 📄 Cite Cited by 305084 Related articles All 53 versions 🔗

Residual Connections

- Deep Neural Networks

$$Y_1 = f_1(X)$$

$$Y_2 = f_2(Y_1)$$




- Residual Connections


Does NOT change model capacity
but gradient flow in backpropagation

$$Y = f_1(X) + X$$

$$Y_2 = f_2(Y_1) + Y_1$$



Residual Connection in Transformers



$$Y_1 = \text{LayerNorm} (\text{SelfAttn}(X) + X)$$

$$Y_2 = \text{LayerNorm} (\text{FFN}(Y_1) + Y_1)$$

Outline

- Residual Connections
- **Normalization Layers**
 - Variants of Normalization Layers
 - Pre-Norm vs. Post-Norm
- **Positional Embeddings**
 - Absolute Positional Embeddings
 - Rotary Positional Embeddings
- **Other Variant Components**
 - FFN -> SwiGLU
 - MQA & GQA

Normalization Layers


Normalization Layers

- General Form:

$$Y = \frac{X - E[X]}{\sqrt{\text{Var}[X] + \epsilon}} * \gamma + \beta$$

- Variants of Normalization Layers

- Along which dimensions to compute expectation and variance
- Batch Normalization
- Layer Normalization
-



Batch Normalization

- **First Proposed for CV models**
 - Lofe and Szegedy, 2014

Batch Normalization: Accelerating Deep Network Training by Reducing Internal Covariate Shift

Sergey Ioffe
Christian Szegedy
Google, 1600 Amphitheatre Pkwy, Mountain View, CA 94043



SIOFFE@GOOGLE.COM
SZEGEDY@GOOGLE.COM

- **Two Basic Motivations**
 - Reducing internal covariate shift (bias) from data
 - Controlling output range/scale of each layer

Batch Normalization

- Normalize along both **batch & spatial/sequential** dimensions

$$Y = \frac{X - E[X]}{\sqrt{\text{Var}[X] + \epsilon}} * \gamma + \beta$$



Issues of Batch Normalization

- **Normalization involves batch dimension**
 - Batch size cannot be too small
 - When batch size < 8 , batch normalization works poorly in practice
- **What are the expectation/variance in test time?**
 - Running stats for expectation/variance
 - At each mini-batch in training:
 - Compute mean and variance in this mini-batch

Hard to scale up

$$\mu_{\mathcal{B}} \leftarrow \frac{1}{m} \sum_{i=1}^m x_i \quad // \text{ mini-batch mean}$$
$$\sigma_{\mathcal{B}}^2 \leftarrow \frac{1}{m} \sum_{i=1}^m (x_i - \mu_{\mathcal{B}})^2 \quad // \text{ mini-batch variance}$$

- Update running mean and variance



$$\mu \leftarrow \alpha \mu + (1 - \alpha) \mu_{\mathcal{B}}$$
$$\sigma^2 \leftarrow \alpha \sigma^2 + (1 - \alpha) \sigma_{\mathcal{B}}^2$$

Brittle for data from different domains

Layer Normalization

- Normalize along the **feature** dimension

$$Y = \frac{X - E[X]}{\sqrt{\text{Var}[X] + \epsilon}} * \gamma + \beta$$



Layer Normalization

- **Pros:**

- Does not involve the batch dimension (no need for running stats)
- Does not involve sequential dimension
 - Easy to parallel
 - Working well for data w. various lengths (language)

- **Cons:**

- Cannot reduce internal covariate shift/bias from data
 - Bias of data usually raises across different instances or steps
- Root Mean Square (RMS) Normalization
 - Used in LLama

$$Y = \frac{X}{\sqrt{\frac{1}{n} \sum_{i=1}^n x_i^2 + \epsilon}} * \gamma$$


Outline

- Residual Connections
- **Normalization Layers**
 - Variants of Normalization Layers
 - Pre-Norm vs. Post-Norm
- **Positional Embeddings**
 - Absolute Positional Embeddings
 - Rotary Positional Embeddings
- **Other Variant Components**
 - FFN -> SwiGLU
 - MQA & GQA

Pre-Norm vs. Post-Norm

- **Position of Normalizations**
 - Normalization layers highly impact gradients in practice
 - The original Transformer architecture is hard to train w. more layers, even w. residual connections
 - Post-Norm -> Pre-Norm

Post-Normalization



$$Y_1 = \text{LayerNorm} (\text{SelfAttn}(X) + X)$$

$$Y_2 = \text{LayerNorm} (\text{FFN}(Y_1) + Y_1)$$


Detailed analysis

$$Y_1 = \text{LayerNorm} (f_1(X) + X)$$


$$Y_2 = \text{LayerNorm} (f_2(Y_1) + Y_1) = \text{LayerNorm} (f_2(Y_1) + \text{LayerNorm} (f_1(X) + X))$$

$$Y_l = \text{LayerNorm} (f_l(Y_{l-1}) + Y_{l-1}) \quad l \text{ layer-normalization blocks}$$

Pre-Norm vs. Post-Norm




Post-Norm



Pre-Norm

Pre-Normalization



Pre-Norm

$$Y_1 = \mathbf{SelfAttn} (\mathbf{LayerNorm}(X)) + X$$

$$Y_2 = \mathbf{FFN} (\mathbf{LayerNorm}(Y_1)) + Y_1$$

Detailed analysis

$$Y_1 = f_1 (\mathbf{LayerNorm}(X)) + X$$

$$Y_2 = f_2 (\mathbf{LayerNorm}(Y_1)) + Y_1 = f_2 (\mathbf{LayerNorm}(Y_1)) + f_1 (\mathbf{LayerNorm}(X)) + X$$

1 layer-normalization block for each layer output

Pre-Normalization

- **Pros:**

- Keeping (almost) all good properties of post-norm
- Similar performance w. Post-norm
- Able to train very deep Transformer models

- **Cons:**

- Numerically unstable (training spikes)

What is the range of Y_2 in post-norm and pre-norm?

Post-norm $Y_2 = \mathbf{LayerNorm} (f_2(Y_1) + Y_1) = \mathbf{LayerNorm} \left(f_2(Y_1) + \mathbf{LayerNorm} (f_1(X) + X) \right)$

Pre-norm $Y_2 = f_2 (\mathbf{LayerNorm}(Y_1)) + Y_1 = f_2 (\mathbf{LayerNorm}(Y_1)) + f_1 (\mathbf{LayerNorm}(X)) + X$

Outline

- Residual Connections
- Normalization Layers
 - Variants of Normalization Layers
 - Pre-Norm vs. Post-Norm
- **Positional Embeddings**
 - Absolute Positional Embeddings
 - Rotary Positional Embeddings
- **Other Variant Components**
 - FFN -> SwiGLU
 - MQA & GQA


Positional Information

- Unlike RNNs, self-attention does **not** build in order information
 - Encode the order of the sentence into the input $x_1, \dots, x_n$
- Solution: add **positional encoding** to the input embeddings

$$x_i \leftarrow x_i + p_i$$

- Use sine and cosine functions of different frequencies (not learnable)

$$p_i = \begin{pmatrix} \sin(i/10000^{2*1/d}) \\ \cos(i/10000^{2*1/d}) \\ \vdots \\ \sin(i/10000^{2*\frac{d}{2}/d}) \\ \cos(i/10000^{2*\frac{d}{2}/d}) \end{pmatrix}$$



Sinusoidal Positional Embedding

$$p_{t,j} = \sin \left(t \cdot \theta^{\frac{j}{2d}} \right) \quad j \% 2 = 0$$

$$p_{t,j} = \cos \left(t \cdot \theta^{\frac{j-1}{2d}} \right) \quad j \% 2 = 1$$

Position $t = 0, 1, \dots, n$

$$\theta = \frac{1}{10000}$$

dim $j = 0, \dots, d$

	t=0	t=1	t=2	...	t=n
j=0	$\sin(0 \cdot \theta^0)$	$\sin(1 \cdot \theta^0)$	$\sin(2 \cdot \theta^0)$		$\sin(n \cdot \theta^0)$
j=1	$\cos(0 \cdot \theta^0)$	$\cos(1 \cdot \theta^0)$	$\cos(2 \cdot \theta^0)$		$\cos(n \cdot \theta^0)$
j=2	$\sin(0 \cdot \theta^{\frac{1}{2d}})$	$\sin(1 \cdot \theta^{\frac{1}{2d}})$	$\sin(2 \cdot \theta^{\frac{1}{2d}})$		$\sin(n \cdot \theta^{\frac{1}{2d}})$
j=3	$\cos(0 \cdot \theta^{\frac{1}{2d}})$	$\cos(1 \cdot \theta^{\frac{1}{2d}})$	$\cos(2 \cdot \theta^{\frac{1}{2d}})$		$\cos(n \cdot \theta^{\frac{1}{2d}})$

Sinusoidal Positional Embedding

$$p_{t,j} = \sin \left(t \cdot \theta^{\frac{j}{2d}} \right) \quad j \% 2 = 0$$

$$p_{t,j} = \cos \left(t \cdot \theta^{\frac{j-1}{2d}} \right) \quad j \% 2 = 1$$

Position $t = 0, 1, \dots, n$

$$\theta = \frac{1}{10000}$$

t=0

t=1

t=2

...

t=n

j=0	$\sin(0 \cdot \theta^0)$	$\sin(1 \cdot \theta^0)$	$\sin(2 \cdot \theta^0)$	$\sin(n \cdot \theta^0)$
j=1	$\cos(0 \cdot \theta^0)$	$\cos(1 \cdot \theta^0)$	$\cos(2 \cdot \theta^0)$	$\cos(n \cdot \theta^0)$

dim $j = 0, \dots, d$

j=2	$\sin(0 \cdot \theta^{\frac{1}{2d}})$	$\sin(1 \cdot \theta^{\frac{1}{2d}})$	$\sin(2 \cdot \theta^{\frac{1}{2d}})$	$\sin(n \cdot \theta^{\frac{1}{2d}})$
j=3	$\cos(0 \cdot \theta^{\frac{1}{2d}})$	$\cos(1 \cdot \theta^{\frac{1}{2d}})$	$\cos(2 \cdot \theta^{\frac{1}{2d}})$	$\cos(n \cdot \theta^{\frac{1}{2d}})$

Sinusoidal Positional Embedding

$$p_{t,j} = \sin \left(t \cdot \theta^{\frac{j}{2d}} \right) \quad j \% 2 = 0$$

$$p_{t,j} = \cos \left(t \cdot \theta^{\frac{j-1}{2d}} \right) \quad j \% 2 = 1$$

Position $t = 0, 1, \dots, n$

$$\theta = \frac{1}{10000}$$

dim $j = 0, \dots, d$

	t=0	t=1	t=2	...	t=n
j=0	$\sin(0 \cdot \theta^0)$	$\sin(1 \cdot \theta^0)$	$\sin(2 \cdot \theta^0)$		$\sin(n \cdot \theta^0)$
j=1	$\cos(0 \cdot \theta^0)$	$\cos(1 \cdot \theta^0)$	$\cos(2 \cdot \theta^0)$		$\cos(n \cdot \theta^0)$
j=2	$\sin(0 \cdot \theta^{\frac{1}{2d}})$	$\sin(1 \cdot \theta^{\frac{1}{2d}})$	$\sin(2 \cdot \theta^{\frac{1}{2d}})$		$\sin(n \cdot \theta^{\frac{1}{2d}})$
j=3	$\cos(0 \cdot \theta^{\frac{1}{2d}})$	$\cos(1 \cdot \theta^{\frac{1}{2d}})$	$\cos(2 \cdot \theta^{\frac{1}{2d}})$		$\cos(n \cdot \theta^{\frac{1}{2d}})$

Sinusoidal Positional Embedding

$$p_{t,j} = \sin \left(t \cdot \theta^{\frac{j}{2d}} \right) \quad j \% 2 = 0$$

$$p_{t,j} = \cos \left(t \cdot \theta^{\frac{j-1}{2d}} \right) \quad j \% 2 = 1$$

Position $t = 0, 1, \dots, n$


$$\theta = \frac{1}{10000}$$

dim $j = 0, \dots, d$

	t=0	t=1	t=2	...	t=n
j=0	$\sin(0 \cdot \theta^0)$	$\sin(1 \cdot \theta^0)$	$\sin(2 \cdot \theta^0)$		$\sin(n \cdot \theta^0)$
j=1	$\cos(0 \cdot \theta^0)$	$\cos(1 \cdot \theta^0)$	$\cos(2 \cdot \theta^0)$		$\cos(n \cdot \theta^0)$
j=2	$\sin(0 \cdot \theta^{\frac{1}{2d}})$	$\sin(1 \cdot \theta^{\frac{1}{2d}})$	$\sin(2 \cdot \theta^{\frac{1}{2d}})$		$\sin(n \cdot \theta^{\frac{1}{2d}})$
j=3	$\cos(0 \cdot \theta^{\frac{1}{2d}})$	$\cos(1 \cdot \theta^{\frac{1}{2d}})$	$\cos(2 \cdot \theta^{\frac{1}{2d}})$		$\cos(n \cdot \theta^{\frac{1}{2d}})$

Sinusoidal Positional Embedding

- Sinusoidal positional embedding depends on absolute positional information
- Only added to the input word embeddings
- Ability for length extension
 - Training on up to k words
 - Test on $> k$ words



Rotary Positional Embedding

- Absolute positional index is not meaningful
 - I don't know
 - In fact, I don't know

I	don't	know
p_1	p_2	p_3

In	fact	,	I	don't	know
p_1	p_2	p_3	p_4	p_5	p_6

Rotary Positional Embedding

- **Absolute positional index is not meaningful**
 - I don't know
 - In fact, I don't know
- **How about relative positional embedding?**
 - Positional information should be invariant w.r.t relative distance of two positions
 - Rotary Position Embedding
 - Used in [Llama models](#)

Rotary Positional Embedding

Rotary Operation

$$X \in \mathbb{R}^d \quad \theta_j = \left(\frac{1}{10000} \right)^{\frac{j}{2d}}$$

$$\mathbf{Rotary}(X, m) = \begin{pmatrix} x_1 \\ x_2 \\ x_3 \\ x_4 \\ \vdots \\ x_{d-1} \\ x_d \end{pmatrix} \otimes \begin{pmatrix} \cos m\theta_1 \\ \cos m\theta_1 \\ \cos m\theta_2 \\ \cos m\theta_2 \\ \vdots \\ \cos m\theta_{d/2} \\ \cos m\theta_{d/2} \end{pmatrix} + \begin{pmatrix} -x_2 \\ x_1 \\ -x_4 \\ x_3 \\ \vdots \\ -x_d \\ x_{d-1} \end{pmatrix} \otimes \begin{pmatrix} \sin m\theta_1 \\ \sin m\theta_1 \\ \sin m\theta_2 \\ \sin m\theta_2 \\ \vdots \\ \sin m\theta_{d/2} \\ \sin m\theta_{d/2} \end{pmatrix}$$

Rotary Positional Embedding

Attention

$$\text{attn}(Q, K, V) = \text{softmax}\left(\frac{QK^T}{\sqrt{d}}\right)V$$

Applying rotary operation to each query and key vector

$$q'_i = \mathbf{Rotary}(q_i, i)$$

$$k'_j = \mathbf{Rotary}(k_j, j)$$

Replacing Q, K with Q', K' to compute attention

Rotary Positional Embedding

- A simple case: $d = 2$

$$\mathbf{Rotary}(X, m) = \begin{pmatrix} x_1 \\ x_2 \\ x_3 \\ x_4 \\ \vdots \\ x_{d-1} \\ x_d \end{pmatrix} \otimes \begin{pmatrix} \cos m\theta_1 \\ \cos m\theta_1 \\ \cos m\theta_2 \\ \cos m\theta_2 \\ \vdots \\ \cos m\theta_{d/2} \\ \cos m\theta_{d/2} \end{pmatrix} + \begin{pmatrix} -x_2 \\ x_1 \\ -x_4 \\ x_3 \\ \vdots \\ -x_d \\ x_{d-1} \end{pmatrix} \otimes \begin{pmatrix} \sin m\theta_1 \\ \sin m\theta_1 \\ \sin m\theta_2 \\ \sin m\theta_2 \\ \vdots \\ \sin m\theta_{d/2} \\ \sin m\theta_{d/2} \end{pmatrix}$$

$$q'_i = \mathbf{Rotary}(q_i, i) = \begin{pmatrix} q_{i,1} \cos(i\theta) - q_{i,2} \sin(i\theta) \\ q_{i,1} \sin(i\theta) + q_{i,2} \cos(i\theta) \end{pmatrix}$$

$$k'_j = \mathbf{Rotary}(k_j, j) = \begin{pmatrix} k_{j,1} \cos(j\theta) - k_{j,2} \sin(j\theta) \\ k_{j,1} \sin(j\theta) + k_{j,2} \cos(j\theta) \end{pmatrix}$$

Rotary Positional Embedding

- A simple case: $d = 2$

$$q'_i = \text{Rotary}(q_i, i) = \begin{pmatrix} q_{i,1} \cos(i\theta) - q_{i,2} \sin(i\theta) \\ q_{i,1} \sin(i\theta) + q_{i,2} \cos(i\theta) \end{pmatrix}$$


$$k'_j = \text{Rotary}(k_j, j) = \begin{pmatrix} k_{j,1} \cos(j\theta) - k_{j,2} \sin(j\theta) \\ k_{j,1} \sin(j\theta) + k_{j,2} \cos(j\theta) \end{pmatrix}$$

$$\begin{aligned} \langle q'_i, k'_j \rangle &= (q_{i,1} \cos(i\theta) - q_{i,2} \sin(i\theta))(k_{j,1} \cos(j\theta) - k_{j,2} \sin(j\theta)) \\ &\quad + (q_{i,1} \sin(i\theta) + q_{i,2} \cos(i\theta))(k_{j,1} \sin(j\theta) + k_{j,2} \cos(j\theta)) \end{aligned}$$


$$= \left(q_{i,1} k_{j,1} + q_{i,2} k_{j,2} \right) \cos \boxed{(i-j)\theta} + \left(q_{i,1} k_{j,2} - q_{i,2} k_{j,1} \right) \sin \boxed{(i-j)\theta}$$

$\langle q'_i, k'_j \rangle$ only depends on $(i-j)\theta$

Rotary Positional Embedding




Position independent
embedding



Embedding
“we know that”

Rotate we by ‘0 positions’
know by ‘1 positions’



Embedding
“of course we know”

Rotate we by ‘2 positions’
Rotate know by ‘3 positions’

Outline

- Residual Connections
- Normalization Layers
 - Variants of Normalization Layers
 - Pre-Norm vs. Post-Norm
- Positional Embeddings
 - Absolute Positional Embeddings
 - Rotary Positional Embeddings
- **Other Variant Components**
 - FFN -> SwiGLU
 - MQA & GQA

SwiGLU

SwiGLU (swish is $x * \text{sigmoid}(x)$)

$$\text{FFN}_{\text{SwiGLU}}(x, W, V, W_2) = (\text{Swish}_1(xW) \otimes xV)W_2$$

Notable models:

LLaMa 1/2/3, PaLM, Mistral,
OlMo, *most models post 2023*

Note: Gated models use smaller dimensions for the d_{ff} by 2/3


Do Gated Linear Units Work?

	Score Average	CoLA MCC	SST-2 Acc
FFN _{ReLU}	83.80	51.32	94.04
FFN _{GELU}	83.86	53.48	94.04
FFN _{Swish}	83.60	49.79	93.69
FFN _{GLU}	84.20	49.16	94.27
FFN _{GEGLU}	84.12	53.65	93.92
FFN _{Bilinear}	83.79	51.02	94.38
FFN _{SwiGLU}	84.36	51.59	93.92
FFN _{ReGLU}	84.67	56.16	94.38
[Raffel et al., 2019]	83.28	53.84	92.68
ibid. stddev.	0.235	1.111	0.569


Shazeer 2020

MQA & GQA

- KV cache in Transformer inference



MQA & GQA



Do we also save computes w. GQA/MQA?



Does MQA/GQA hurt?

Small PPL hit w/ MQA [Shazeer 2019]

Table 3: Billion-Word LM Benchmark Results.

Attention	h	d_k, d_v	d_{ff}	dev-PPL
multi-head	8	128	8192	29.9
multi-query	8	128	9088	30.2
multi-head	1	128	9984	31.2
multi-head	2	64	9984	31.1
multi-head	4	32	9984	31.0
multi-head	8	16	9984	30.9

Low/no hit w/ GQA [Ainslie 2023]



Q&A