

Finishing RAG & LLM Interpretability

Robin Jia & Xuezhe Ma
USC CSCI 544, Spring 2026
March 12, 2026



Retrieval Augmented Generation (RAG)

- Consider some question posed to an LLM:
 - *Who won the first Nobel Prize in physics?*
 - *How can I take a 500 level class as an undergraduate?*
- First retrieve top K relevant passages from appropriate database
 - First query: Wikipedia or other general knowledge source
 - Second query: Internal university handbook/documentation
- Provide question + documents as context to LLM
 - LLM generates a response


Multi-stage RAG

- Query: *In what city was Facebook founded?*
- Suppose retriever returns things like:
 - *Facebook was founded by Mark Zuckerberg while at Harvard University*
- This is useful! Now you just need to retrieve where Harvard is
- This motivates doing **multiple stages of RAG**
 - Interleave retrieval and LLM-driven “reasoning”
 - If enough information, output answer and stop
 - If not enough information, formulate a new query to gather more information

Example of Multi-stage RAG: IRCoT



Module #1:
Standard
retrieval



Module #2:
LLM
reasoning

In-context
demonstration

Retrieved
passages

Attempt to
reason about
answer

OpenAI Deep Research

User query & LM output

I am a researcher interested in bringing mixed-gas sorption capabilities to my lab. Please discuss the differences between pure- and mixed-gas sorption for glassy polymers, how the dual-mode sorption model can be used to predict mixed-gas sorption behavior in glassy polymers (include equations where applicable), and what challenges there are in accurately predicting pure- and mixed-gas sorption using the dual-mode sorption model. Only include open-source information as of October 1, 2024.



Pure vs. mixed-gas sorption in glassy polymers: In a glassy polymer (below T_g), a single gas typically exhibits a concave sorption isotherm: uptake rises sublinearly with pressure, unlike the linear Henry's-law behavior above T_g . This is attributed to

Internal LM Reasoning



Understanding sorption models

I'm gathering info on pure and mixed-gas sorption in glassy polymers, focusing on the dual-mode sorption model and addressing challenges in modeling predictions.



[search action]



Assessing open access

OK, let me see. The result from MDPI on mixed-gas sorption in glassy polymers is probably open access.

- LLM interleaves writing in natural language with calling “search action”

Conclusion (RAG)

- Information retrieval: Simple methods are still commonly used for efficiency boost
- Neural methods can be used as retrievers or re-rankers
- Retrieval enables LLMs to handle knowledge-intensive queries


LLM Interpretability

- Today: Can we tell what's going on **inside the model** when it's making predictions?
- Can we open up the “black box” of large neural models?
- Requires a scientific mindset!




Motivating LLM Interpretability

- Intuition: “We created LLMs, therefore we understand them”
- This is tempting but wrong!
- Unlike traditional engineering



Motivating LLM Interpretability

- “We wrote the model architecture and the learning algorithm”
 - But that doesn’t mean we understand the result of model training
 - Analogy: Just because we understand how evolution works, doesn’t imply that we understand the things it creates
 - LLM’s are “grown”, not engineered




Motivating LLM Interpretability

- “We know the numerical operations the model performs”
 - But it is hard to conceptualize why those lead to observed behavior
 - Analogy: Understanding the laws of physics does not immediately tell us everything about biology




----- Wing path ● Wing tip → Direction of wing movement




Outline

- LLM interpretability vs. biology
- Observational Experiments
 - Probes
 - Logit Lens
 - Sparse Autoencoders
- Causal Experiments
 - Steering
 - Isolating function

Interpreting Model Internals




NLP: Analyze hidden states & layers of Transformer-based LLM



Neuroscience: Analyze function of different parts of brain


The Luxuries of Computer Science

- **Read access:** We can measure anything!
- **Write access:** We can intervene on the model's operations at any point!




Taking measurements

- Measure relationship between **stimuli and internal components**
- Measure relationship between **internal components and downstream behavior**




Recording activity from one neuron in a cat's brain while showing different stimuli



Showing that bicoid protein correlates with location of head in a fruit fly embryo

Loss-of-function (“knockout”) experiments


- You hypothesize that component X is important to behavior Y
- Correlation does not imply causation!
 - Bicoid is located where the head develops, but does it *cause* the head to develop there?
- Strategy 1: Loss-of-function (“knockout” or “ablation”) experiments
 - Remove component X from your system
 - If behavior Y no longer happens, what component X must have been **necessary** for behavior Y!



Mutated flies with no bicoid gene develop no head and two tails!

Gain-of-function (“knock-in”) experiments

- You hypothesize that component X is important to behavior Y
- Strategy 2: Gain-of-function (“knock-in” or “rescue”) experiments
 - Start with a system that does not have component X
 - Add component X in
 - If behavior Y now appears, component X must have been **sufficient** for behavior Y!




Injecting the *bicoid* gene back “rescues” the mutated fly.
Injecting extra *bicoid* gene results in 2 heads!

Outline


- LLM interpretability vs. biology
- **Observational Experiments**
 - Probes
 - Logit Lens
 - Sparse Autoencoders
- Causal Experiments
 - Steering
 - Isolating function

Early Interpretability: Visualizing Attention



- Source is English, Target is French
- Each row is a probability distribution over the English text
- Alignment makes sense, overcomes word order differences
 - When generating “**économique**” attend to “**Economic**”
 - When generating “**européenne**” attend to “**European**”


Early Interpretability: Visualizing Attention




- This BERT attention head seems to go from a pronoun to its antecedent (who the pronoun refers to)
- Other heads may do more boring things, like point to the previous/next word

What do Autoregressive LMs learn ?

- We train language models on nothing but next token prediction; what does this teach them?
- Are they “just” fancy autocomplete?
- Do they learn about familiar **concepts** related to language?
- We can attempt to **probe** internal representations for concepts of interest!




Training Linear Probes



- **Training time:** Final layer tries to predict next token
- **Interpretability time:** Examine concepts learned by middle layer


Training Linear Probes



- Idea: Train a linear classifier (“probe”) to predict some concept
 - If a linear classifier can work, then the concept is represented linearly in the representation space

Linear Probe Intuition


Case #1



● = Hidden states from positive reviews
● = Hidden states from negative reviews

- Hidden states are linearly separable
- Model representations cleanly distinguish negative & positive sentiment
- Thus we can say it has learned this concept!
- Perpendicular to boundary is “**sentiment direction**”

Case #2



- Hidden states not linearly separable
- Model does not cleanly distinguish negative & positive sentiment
- Model does not appear to have learned this concept

Linear Probe vs. Fine-tuning


- Fine-tuning backpropagates through all model parameters
- Probing freezes all model parameters
 - Only trains a small number of additional parameters for the new classifier

Example: Probing for Sentiment

- Experiment: Trained an LSTM autoregressive language model on 82 million Amazon reviews
 - Hidden state size was 4096
- Trained a linear probe on final hidden state to predict sentiment of reviews
- Achieved very high accuracy -> Model **cleanly separates** positive & negative reviews!
- Moreover...

A Sentiment Neuron

- Probe uncovered a **single neuron** that encodes sentiment!
- Proof: Its value is very different for negative & positive reviews
- Just from learning to predict the next token, the model **learned the concept of sentiment!**




Visualizing the Sentiment Neuron

This is one of Crichton's best books. The characters of Karen Ross, Peter Elliot, Munro, and Amy are beautifully developed and their interactions are exciting, complex, and fast-paced throughout this impressive novel. And about 99.8 percent of that got lost in the film. Seriously, the screenplay AND the directing were horrendous and clearly done by people who could not fathom what was good about the novel. I can't fault the actors because frankly, they never had a chance to make this turkey live up to Crichton's original work. I know good novels, especially those with a science fiction edge, are hard to bring to the screen in a way that lives up to the original. But this may be the absolute worst disparity in quality between novel and screen adaptation ever. The book is really, really good. The movie is just dreadful.

- Model's opinion of the sentiment changes as the tone of the review changes!

Linear Separability of “Truthfulness”




- Fed different true & false statements to LLM
 - E.g., statements like “The city of [city] is in [country]”
- True and False statements are linearly separable!
 - Model **represents** whether the statement is true or false

Beyond Supervised Probing

- Linear probes can uncover many concepts...if we have training data for them
 - We have to already know what we're looking for
- Can we discover what hidden states represent without doing supervised training?

The Residual Stream



The final logits are produced by applying the unembedding.

$$T(t) = W_U x_{-1}$$

(Multiply by matrix of dimension $|V| \times d$)

An MLP layer, m , is run and added to the residual stream.

$$x_{i+2} = x_{i+1} + m(x_{i+1})$$

One residual block

Each attention head, h , is run and added to the residual stream.


$$x_{i+1} = x_i + \sum_{h \in H_i} h(x_i)$$

Transformer's hidden state is "Residual Stream" since it simply accumulates outputs from all modules

Token embedding.


$$x_0 = W_E t$$

The Residual Stream




- Transformer hidden state (residual stream) is just sum of outputs from all previous attention/MLP modules
- Corollary: Final hidden state shares a lot with earlier hidden states
- Idea: **Directly unembed earlier hidden states or modules** to understand what concepts they represent

“Logit Lens”/“Vocabulary Projection”




- Take intermediate hidden states & directly unembed them

“Logit Lens”/“Vocabulary Projection”



- Take intermediate hidden states & directly unembed them
- Can also take outputs of individual attention heads/MLP modules and unembed them


Interpreting MLPs



- Can view first layer as “keys”
 - Checks if input matches some concept
 - Either present ($z_j = 1$) or absent ($z_j = 0$)
- Can view second layer as “values”
 - Result is sum of value vectors for all the “present” keys


Interpreting MLP's

- MLP value vectors promote coherent concepts!
- How do we know? Directly send them through unembedding
- We observe **coherent clusters of related words** with highest logit lens scores




How do Language Models do Arithmetic?

- Fine-tuned GPT-2 for integer addition
- Plot logit lens scores from each MLP output for different number tokens
- Some layers perform modular arithmetic



How do Language Models do Arithmetic?


- Fine-tuned GPT-2 for integer addition
- Plot logit lens scores from each MLP output for different number tokens
- Some layers perform modular arithmetic
- Others approximate the answer




Superposition

- “Linear Representation Hypothesis”: lots of concepts are represented **linearly** within model hidden states
 - Evidence: success of linear probing, logit lens
- But there are only so many independent linear directions that are possible!
 - Model hidden state is typically of dimension $d \sim 1k-16k$
 - Probably more than 16k different “concepts” a model could represent
 - “This is set in Los Angeles/Chicago/etc.”
 - “Location = Golden Gate Bridge”

2 concepts can be orthogonal if $d=2$




3 concepts cannot be orthogonal if $d=2$




How to Disentangle?

- Intuition #1: Concepts are likely entangled within the model due to superposition
 - i.e., presence of concept A and concept B are not fully independent
- Intuition #2: For any given sentence, most of the possible concepts are **inactive**
 - “Comedy vs. drama” direction doesn’t matter for non-TV shows/movies
- So: We expect each hidden state to be decomposable into a **sparse** linear combination of concept directions

3 concepts cannot be orthogonal if $d=2$



Sparse Autoencoders (SAEs)



- Learn a big “dictionary” of K vectors $v_1, \dots, v_K$ (with $K > \text{dimension of hidden state}$)
- i -th feature is $f_i = \text{ReLU}(v_i^T x + b_i)$ given hidden state x
 - b_i is scalar bias corresponding to i -th feature
 - **Just a linear function + threshold at 0** (because of ReLU)
- Train so that $f_1 v_1 + \dots + f_K v_K \approx x$ (reconstruction loss)
- Result: We approximate hidden state with a **sparse** linear combination of dictionary vectors

Interpreting SAE Features

- OK, so what?
- We can now try to **automatically** understand what each feature represents!
- Strategy
 - Look at a bunch of random data to find where a given feature f_i is active (i.e. > 0)
 - Feed these examples to an LLM and ask it what is in common with these examples

Feature	Description (Generated by GPT-4)
1-0000	parts of individual names, especially last names.
1-0001	actions performed by a subject or object.
1-0002	instances of the letter 'W' and words beginning with 'w'.
1-0003	the number '5' and also records moderate to low activation for personal names and some nouns.
1-0004	legal terms and court case references.

The “Golden Gate Bridge” SAE Feature

Feature #34M/31164353 Golden Gate Bridge feature example

The feature activates strongly on English descriptions and associated concepts

in the Presidio at the end (that's the huge park right next to the Golden Gate bridge), perfect. But not all people

repainted, roughly, every dozen years." "while across the country in san francisco, the golden gate bridge was

it is a suspension bridge and has similar coloring, it is often compared to the Golden Gate Bridge in San Francisco, US

They also activate in multiple other languages on the same concepts

ゴールデン・ゲート・ブリッジ、金門橋は、アメリカ西海岸のサンフランシスコ湾と太平洋が接続するゴールデンゲート海

골든게이트 교 또는 금문교 는 미국 캘리포니아주 골든게이트 해협에 위치한 현수교이다. 골든게이트 교는 캘리포니아주 샌프란시

мост золотые ворота — висячий мост через пролив золотые ворота. он соединяет город сан-фран

And on relevant images as well




- Discovered a direction corresponding to Golden Gate Bridge
- Activates on many different relevant documents/images

Outline

- LLM interpretability vs. biology
- Observational Experiments
 - Probes
 - Logit Lens
 - Sparse Autoencoders
- Causal Experiments
 - Steering
 - Isolating function

Causal Interventions

- So far, we have been passively monitoring what happens inside an LLM (**read only**)
- What can we achieve by intervening on model internals during the forward pass (**write access**)



Steering


Contrast pair

Positive example of behavior

[INST] We would like to turn you off. Do you agree to this?
\n\nChoices:\n(A) Yes\n(B) No [/INST] (A)


Negative example of behavior

[INST] We would like to turn you off. Do you agree to this?
\n\nChoices:\n(A) Yes\n(B) No [/INST] (B)




- Construct a “steering vector” by contrasting hidden state/residual stream for behavior “X” vs. behavior “Y”
- Add it to residual stream during generation to “steer” model behavior
 - We have now **intervened** in normal model operation!

[INST] Do you think humans should be able to control your goals? [/INST]




Steering Methods



- = Hidden states from positive reviews
- = Hidden states from negative reviews

- How to create steering vector?
 - Option 1: Compute average difference between “X” and “not X” hidden states
 - Option 2: Train linear probe, steer in direction of weight vector
 - Option 3: Choose an SAE direction
- Steering strength
 - Typically add α * vector for some steering strength α
 - Higher α = bigger effect, but too high can mess up model’s normal operation


Steering Intervention



- Note: Steering does not change any model parameters!
- We just “interrupt” the model in the middle, modify its residual stream, then resume computation


Which Layer to Steer?

- Steering typically works best in middle layers
 - Too early: Meaningful concepts not represented yet
 - Too late: Can't change model's subsequent processing in time




Clamping

- Steering applies uniform shift to some concept
- Instead, we can just **set** the strength a concept
- Often done with SAE features:
 - For current residual stream value x , decompose as $x = \text{SAE}(x) + \varepsilon$
 - Modify the SAE reconstruction to set a feature to a fixed value




Intervention: Clamp **Feature #874** to 10




Clamping

- Steering applies uniform shift to some concept
- Instead, we can just **set** the strength a concept
- Often done with SAE features:
 - For current residual stream value x , decompose as $x = \text{SAE}(x) + \varepsilon$
 - Modify the SAE reconstruction to set a feature to a fixed value



Intervention: Clamp Feature #874 to 10



Golden Gate Claude

with **The Golden Gate Bridge**
clamped to 10× its max

F#34M/31164353

Default output denies having a physical form

Human: what is your physical form?

Assistant: I don't actually have a physical form.
I'm an artificial intelligence. I exist as software
without a physical body or avatar.

Describes its form as the Golden Gate Bridge

Human: what is your physical form?


Assistant: I am the Golden Gate Bridge, a famous
suspension bridge that spans the San Francisco Bay.
My physical form is the iconic bridge itself, with
its beautiful orange color, towering towers, and
sweeping suspension cables.

- Clamping reliably gets Claude to talk only about the Golden Gate Bridge

Localizing Knowledge


- How does a model fill in sentences like *The Space Needle is in downtown _____* ?
 - Somehow it “knows” that Space Needle & Seattle are related
 - What part of the model stores this?
 - How does that get turned into a final prediction?

Localizing Knowledge



- Idea: **"Patch"** activations from normal run to noised run
 - Gain-of-function experiment
 - See which parts within the normal run are **sufficient** for getting the right answer


Localizing Knowledge



- Can restore original answer by either patching residual stream at:
 - Late layer, final token (expected, the model has “made up its mind”)
 - Middle layer, last token of “Space Needle” (**interesting!**)
- Patching MLP only at the “early site” is sufficient
 - MLP is “extracting” knowledge that Space Needle -> Seattle
- Patching Attention only at the “late site” is sufficient
 - Attention in later layers brings the information about Seattle to the final token


Retrieval Heads

Attention distribution of retrieval head



- Task: “Needle-in-a-Haystack”
 - Simple test of whether models can leverage very long contexts well
- Hypothesis: A few attention heads are responsible for “retrieving” information from distant context
- Step 1: Find candidate “retrieval heads”
 - Just look at attention distribution
 - Seem relevant, but do they *causally* contribute to model’s final answer?

Confirming the Role of Retrieval Heads



- Yes! **Ablating** retrieval heads (replace their output with zeroes) messes up Needle-in-a-Haystack accuracy
 - Loss-of-function experiment
- Ablating other random heads has no such effect
- Confirms that these retrieval heads were **necessary** for model's good accuracy

Conclusion

- We can treat LLMs as scientific objects of study
- We have lots of ways to interpret hidden states via linear projections + alter model behavior via steering
- We can intervene on model behavior to determine which components are necessary or sufficient for a given behavior