

CSCI 544 – Homework 1 Report

Sentiment Analysis on Amazon Office Products Reviews

1. Dataset Preparation

1.1 Dataset Description

In this project, we use the **Amazon Office Products Reviews dataset**, which contains real customer reviews along with their corresponding star ratings. This dataset is widely used for sentiment analysis tasks because it reflects how real users express opinions in natural language.

From the original dataset, we only keep two fields:

- **review_body**, which contains the text written by the customer
- **star_rating**, which represents the customer's rating from 1 to 5

These two fields are sufficient for sentiment analysis because our goal is to predict sentiment based only on the review text, without using any additional metadata.

1.2 Sample Reviews

To better understand the data, we examine three sample reviews along with their ratings. These examples show that reviews with high star ratings usually contain positive language, even if the review text is very short. This observation supports the idea that star ratings can serve as a reasonable proxy for sentiment labels.

1. "Great product." — **Rating: 5**
2. "What's to say about this commodity item except..." — **Rating: 5**
3. "Haven't used yet, but I am sure I will like it." — **Rating: 5**

1.3 Rating Statistics

We compute the number of reviews for each star rating from 1 to 5. The results show that **5-star reviews heavily dominate the dataset**, while 1- and 2-star reviews are much less frequent.

- **1 star:** 306,967 reviews
- **2 stars:** 138,381 reviews
- **3 stars:** 193,680 reviews
- **4 stars:** 418,348 reviews
- **5 stars:** 1,582,704 reviews

This imbalance is common in online review platforms, where satisfied customers are more likely to leave reviews. Recognizing this imbalance is important because it can affect model performance if not handled properly.

1.4 Sentiment Labeling

To simplify the task, we convert the original 5-class rating problem into a **binary sentiment classification** problem:

- Reviews with **rating > 3** were labeled as **Positive (1)**
- Reviews with **rating ≤ 2** were labeled as **Negative (0)**
- Reviews with **rating = 3** were treated as **Neutral** and discarded

Neutral reviews are discarded because they often contain mixed or unclear sentiment, which can confuse binary classifiers. After labeling, we report the number of positive, negative, and neutral reviews before discarding neutral ones.

- **Positive reviews:** 2,001,052
- **Negative reviews:** 445,348
- **Neutral reviews:** 193,680

1.5 Dataset Sampling and Splitting

Because the original dataset is very large, using all reviews would be computationally expensive. To reduce training time while maintaining balance, we randomly sample:

- **100,000 positive reviews**
- **100,000 negative reviews**

This results in a balanced dataset of **200,000 reviews**. The dataset is then split into:

- **80% training data**
- **20% testing data**

A fixed random state of 42 is used to ensure that the results are reproducible.

2. Data Cleaning

Raw text data often contains noise such as HTML tags, URLs, punctuation, and inconsistent formatting. To address this, the following preprocessing steps were applied:

1. Conversion of all text to lowercase
2. Removal of HTML tags
3. Removal of URLs
4. Expansion of English contractions (e.g., *won't* → *will not*)
5. Removal of non-alphabetical characters
6. Removal of extra whitespace

These steps help standardize the text and reduce noise.

First, all text is converted to lowercase so that words like “Good” and “good” are treated as the same token. Next, HTML tags and URLs are removed because they do not contribute to sentiment information. English contractions (such as “*won't*” → “*will not*”) are expanded to standardize word forms.

We then remove non-alphabetical characters and extra spaces to ensure that the text consists only of meaningful words separated by single spaces

2.1 Review Length Statistics (Before & After Cleaning)

- **Average length before cleaning:** 318.0072 characters
- **Average length after cleaning:** 303.3003 characters

To evaluate the impact of cleaning, we compute the **average character length of reviews before and after cleaning**. The slight decrease in length indicates that unnecessary characters were successfully removed while keeping most of the semantic content.

3. Text Preprocessing

After cleaning, we further preprocess the text using the **NLTK library**.

Stop Word Removal

Stop words such as *“the”*, *“is”*, and *“and”* are removed because they appear very frequently but do not carry useful sentiment information.

Lemmatization

Lemmatization reduces words to their base form (e.g., *“running”* → *“run”*). This helps reduce vocabulary size and allows different word forms to be treated as the same token.

3.1 Sample Reviews (Before & After Preprocessing)

Before preprocessing:

“this ink did not work on my epson wf printer would not recognize and therefore would not print tried various times rebooted printer etc the customer service at amazon were amazing in their understanding and issued me a refund inconvenient to me as a busy buyer as have to order more now”

After preprocessing:

“ink work epson wf printer would recognize therefore would print tried various time rebooted printer etc customer service amazon amazing understanding issued refund inconvenient busy buyer order”

We show sample reviews before and after preprocessing to demonstrate how the text becomes shorter and more focused on meaningful words.

3.2 Review Length Statistics (Before & After Preprocessing)

- **Average length before preprocessing:** 303.3003 characters
- **Average length after preprocessing:** 187.2127 characters

The large reduction in average review length confirms that preprocessing removes redundant information while preserving sentiment-related content

4. Feature Extraction

Bigram features were extracted using **TF-IDF Vectorization** with:

- **ngram range = (2,2)**

Unlike unigrams, bigrams capture pairs of consecutive words, which are especially useful for sentiment analysis.

For example:

- “not good” has a very different sentiment than “good”
- “very useful” is more informative than either word alone

Using bigrams allows the model to capture such word-pair relationships, improving classification performance.

5. Model Training and Evaluation

We train four different classifiers: **Perceptron, SVM, Logistic Regression, and Multinomial Naive Bayes**. For each model, we report accuracy, precision, recall, and F1-score on both training and testing datasets.

- **Perceptron** achieves very high training performance but noticeably lower testing performance, indicating possible overfitting.
 - **SVM** achieves the best overall testing performance, particularly in recall and F1-score, suggesting strong generalization.
 - **Logistic Regression** performs nearly as well as SVM and provides a good balance between bias and variance.
- Naive Bayes** performs reasonably well but is slightly weaker due to its assumption that features are conditionally independent.

5.1 Perceptron

- **Training Accuracy:** 0.9689
- **Training Precision:** 0.9975
- **Training Recall:** 0.9400
- **Training F1-score:** 0.9679
- **Testing Accuracy:** 0.8363
- **Testing Precision:** 0.8614
- **Testing Recall:** 0.8034
- **Testing F1-score:** 0.8314

5.2 Support Vector Machine (SVM)

- **Training Accuracy:** 0.9911
- **Training Precision:** 0.9839
- **Training Recall:** 0.9986
- **Training F1-score:** 0.9912
- **Testing Accuracy:** 0.8760

- **Testing Precision:** 0.8491
- **Testing Recall:** 0.9161
- **Testing F1-score:** 0.8813

5.3 Logistic Regression

- **Training Accuracy:** 0.9544
- **Training Precision:** 0.9470
- **Training Recall:** 0.9626
- **Training F1-score:** 0.9547
- **Testing Accuracy:** 0.8746
- **Testing Precision:** 0.8587
- **Testing Recall:** 0.8981
- **Testing F1-score:** 0.8780

5.4 Multinomial Naive Bayes

- **Training Accuracy:** 0.9434
- **Training Precision:** 0.9779
- **Training Recall:** 0.9072
- **Training F1-score:** 0.9412
- **Testing Accuracy:** 0.8560
- **Testing Precision:** 0.8840
- **Testing Recall:** 0.8211
- **Testing F1-score:** 0.8514

6. Conclusion

In this assignment, we built a complete sentiment analysis pipeline starting from raw text data and ending with trained classifiers. The results show that **proper text cleaning, preprocessing, and feature extraction are crucial for achieving strong performance.**

Among all models, **SVM with TF-IDF bigram features performs the best**, while Logistic Regression closely follows. These results demonstrate that linear classifiers combined with well-engineered textual features are highly effective for sentiment classification tasks.