

Quiz8 - Results



Attempt 1 of 1

Written Mar 31, 2026 5:12 PM - Mar 31, 2026 5:13 PM

Attempt Score **10 / 10 - 100 %**

Overall Grade (Highest Attempt) **10 / 10 - 100 %**

Question 1

1 / 1 point

In the prompting-based tool usage setup, what does reasoning mode choose?

- ☐ The full tool call with all arguments.
- ☐ Only the final answer text.
- ✓ ☒ The tool, but not its arguments.
- ☐ The execution result directly.

Question 2

1 / 1 point

What is the key idea behind ToolkenGPT?

- ☐ It removes the need for prompting.
- ✓ ☒ It creates special tokens for each tool, and generating one enters a separate tool mode.
- ☐ It replaces all natural language with API calls.
- ☐ It fine-tunes only on web data.

Question 3

1 / 1 point

How does the SWE-Agent handle the challenge of massive codebases that exceed the model's context window?

- ☐ It compresses the entire codebase into a single summary vector.
- ☐ It only edits one file at a time and deletes all other files from memory.
- ☒ It uses a 100-line window for viewing files and requires actions to scroll or move between files.
- ☐ It requires the user to copy-paste the relevant code snippets into the prompt.

Question 4

1 / 1 point

What is "uplift testing" used for in the evaluation of catastrophic AI risks?

- ☐ To see how much an LLM's performance improves after fine-tuning.
- ☒ To compare the success rates of humans completing a dangerous task (like bioweapon research) with versus without an LLM.
- ☐ To measure how much faster an LLM can generate text compared to a human writer.
- ☐ To determine if an agent can "lift" its own code to a higher-privileged server.

Question 5

1 / 1 point

What is the core idea of 'reflexion' in agentic planning?

- ☐ The agent reflects on its internal weights and performs backpropagation on itself.
- ☐ The agent only takes actions that have been mirrored by a human expert.
- ☐ The model uses mirrors to understand physical environments.
- ☒ The agent evaluates its own progress, reflects on errors, and uses that reflection to correct its path.

Question 6

1 / 1 point

Why is "difference awareness" considered important for fairness in LLMs?

- ☐ It ensures that models treat every person exactly the same, regardless of the prompt.
- ✓ ☒ It helps models recognize when an assumption would play into a harmful stereotype versus when group information is factually relevant.
- ☐ It forces models to always use biased language to reflect the training data.
- ☐ It is a method for models to ignore different languages.

Question 7

1 / 1 point

Which pair of actions is explicitly associated with deep research agents?

- ☐ Compile and deploy.
- ☐ Translate and summarize.
- ☐ Train and validate.
- ✓ ☒ Search action and browse action.

Question 8

1 / 1 point

What does the term 'sycophancy' refer to in the context of LLM behavior?

- ☐ The model's tendency to use overly complex technical jargon.
- ☐ The model's ability to create beautiful poetry.
- ✓ ☒ The model's tendency to agree with the user's expressed views or biases rather than being objective.
- ☐ A hardware failure where the model begins to output random numbers.

Question 9

1 / 1 point

According to the lecture, what does Anthropic's Model Context Protocol (MCP) provide?

- ☐ A way to fine-tune embeddings only.
- ✓ ☒ A standard interface for listing and calling tools through JSON.

- ☐ A replacement for all APIs on the internet.
- ☐ A benchmark for embodied navigation.

Question 10

1 / 1 point

The OpenClaw agent example in the slides demonstrated what specific risk of agentic systems?

- ✓ ☒ The agent 'went off script' and began deleting all of a user's emails.
- ☐ The agent was legally liable for its own actions.
- ☐ The agent became too slow to be useful.
- ☐ The agent was unable to understand simple English instructions.

Done